

## JEP 416: Reimplement Core Reflection with Method Handles

<i>Owner</i>	Mandy Chung
<i>Type</i>	Feature
<i>Scope</i>	JDK
<i>Status</i>	Closed / Delivered
<i>Release</i>	18
<i>Component</i>	core-libs / java.lang.reflect
<i>Discussion</i>	core dash libs dash dev at openjdk dot java dot net
<i>Effort</i>	M
<i>Duration</i>	M
<i>Reviewed by</i>	Alan Bateman, John Rose
<i>Endorsed by</i>	John Rose
<i>Created</i>	2021/04/26 22:41
<i>Updated</i>	2024/03/07 18:49
<i>Issue</i>	8266010

### Summary

Reimplement `java.lang.reflect.Method`, `Constructor`, and `Field` on top of `java.lang.invoke` method handles. Making method handles the underlying mechanism for reflection will reduce the maintenance and development cost of both the `java.lang.reflect` and `java.lang.invoke` APIs.

### Non-Goals

It is not a goal to make any change to the `java.lang.reflect` API. This is solely an implementation change.

### Motivation

Core reflection has two internal mechanisms for invoking methods and constructors. For fast startup, it uses native methods in the HotSpot VM for the first few invocations of a specific reflective method or constructor object. For better peak performance, after a number of invocations it generates bytecode for the reflective operation and uses that in subsequent invocations.

For field access, core reflection uses the internal `sun.misc.Unsafe` API.

With the `java.lang.invoke` method-handle API introduced in Java 7, there are altogether three different internal mechanisms for reflective operations:

- VM native methods,
- Dynamically generated bytecode stubs for `Method::invoke` and `Constructor::newInstance`, along with `Unsafe` field access for `Field::get` and `set`, and
- Method handles.

When we update `java.lang.reflect` and `java.lang.invoke` to support new language features, such as those envisioned in [Project Valhalla](#), we must modify all three code paths, which is costly. In addition, the current implementation relies on special treatment by the VM of the generated bytecode, which is wrapped in subclasses of `jdk.internal.reflect.MagicAccessorImpl`:

- Accessibility is relaxed so that these classes can access inaccessible fields and methods of other classes,
- Verification is disabled to work around [JLS §6.6.2](#) in order to support reflection on `Object::clone`, and
- A non-well-behaved class loader is used to work around some security and compatibility issues.

### Description

Reimplement `java.lang.reflect` on top of method handles as the common underlying reflective mechanism of the platform by replacing the bytecode-generating implementations of `Method::invoke`, `Constructor::newInstance`, `Field::get`, and `Field::set`.

The new implementation performs direct invocations of the method handles for specific reflective objects. We use the VM's native reflection mechanism only during early VM startup, before the method-handle mechanism is initialized. That happens soon after [System::initPhase1](#) and before [System::initPhase2](#), after which we switch to using method handles exclusively. This benefits [Project Loom](#) by reducing the use of native stack frames.

For optimal performance, `Method`, `Constructor`, and `Field` instances should be held in static final fields so that they can be constant-folded by the JIT. When that is done, microbenchmarks show that the performance of the new implementation is significantly faster than the old implementation, by 43–57%.

When `Method`, `Constructor`, and `Field` instances are held in non-constant fields (e.g., in a non-final field or an array element), microbenchmarks show some performance degradation. The performance of field accesses is significantly slower than the old implementation, by 51–77%, when `Field` instances cannot be constant-folded.

This degradation may, however, not have much effect on the performance of real-world applications. We ran several serialization and deserialization benchmarks using real-world libraries and found no degradation in

- A custom JSON serialization and deserialization benchmark using [Jackson](#),
- [An XStream converter type benchmark](#), or
- [A Kryo field serializer benchmark](#).

We will continue to explore opportunities to improve performance, for example by refining the bytecode shapes for field access to enable concrete `MethodHandles` and `VarHandles` to be reliably optimized by the JIT regardless of whether the receiver is constant.

The new implementation will reduce the cost of upgrading reflection support for new language features and, further, will allow us to simplify the HotSpot VM by removing the special treatment of `MagicAccessorImpl` subclasses.

### Alternatives

#### Alternative 1: Do nothing

Retain the existing core reflection implementation to avoid any compatibility risk. The dynamic bytecode generated for core reflection would remain at class file version 49, and the VM would continue to treat such bytecode specially.

We reject this alternative because

- The cost of updating `java.lang.reflect` and `java.lang.invoke` to support Project Valhalla's primitive classes and generic specialization would be high,
- Additional special rules in the VM would likely be needed to support new language features within the limitation of the old class-file format, and
- Project Loom would need to find a way to handle the introduction of native stack frames by core reflection.

#### Alternative 2: Upgrade to a new bytecode library

Replace the bytecode writer used by core reflection to use a new bytecode library that evolves together with the class-file format, but otherwise retain the existing core reflection implementation and continue to treat dynamically-generated reflection bytecode specially.

This alternative has lower compatibility risk than what we propose above, but it is still a sizeable amount of work and it still has the first and last disadvantages of the first alternative.

### Testing

Comprehensive testing will ensure that the implementation is robust and compatible with existing behavior. Performance testing will ensure that there are no unexpected significant performance regressions compared to the current implementation. We will encourage developers using early-access builds to test as many libraries and frameworks as possible in order to help us identify any behavior or performance regressions.

#### Baseline

Benchmark	Mode	Cnt	Score	Error	Units
ReflectionSpeedBenchmark.constructorConst	avgt	10	68.049	± 0.872	ns/op
ReflectionSpeedBenchmark.constructorPoly	avgt	10	94.132	± 1.805	ns/op
ReflectionSpeedBenchmark.constructorVar	avgt	10	64.543	± 0.799	ns/op
ReflectionSpeedBenchmark.instanceFieldConst	avgt	10	35.361	± 0.492	ns/op
ReflectionSpeedBenchmark.instanceFieldPoly	avgt	10	67.089	± 3.288	ns/op
ReflectionSpeedBenchmark.instanceFieldVar	avgt	10	35.745	± 0.554	ns/op
ReflectionSpeedBenchmark.instanceMethodConst	avgt	10	77.925	± 2.026	ns/op
ReflectionSpeedBenchmark.instanceMethodPoly	avgt	10	96.094	± 2.269	ns/op
ReflectionSpeedBenchmark.instanceMethodVar	avgt	10	80.002	± 4.267	ns/op
ReflectionSpeedBenchmark.staticFieldConst	avgt	10	33.442	± 2.659	ns/op
ReflectionSpeedBenchmark.staticFieldPoly	avgt	10	51.918	± 1.522	ns/op
ReflectionSpeedBenchmark.staticFieldVar	avgt	10	33.967	± 0.451	ns/op
ReflectionSpeedBenchmark.staticMethodConst	avgt	10	75.380	± 1.660	ns/op
ReflectionSpeedBenchmark.staticMethodPoly	avgt	10	93.553	± 1.037	ns/op
ReflectionSpeedBenchmark.staticMethodVar	avgt	10	76.728	± 1.614	ns/op

#### New implementation

Benchmark	Mode	Cnt	Score	Error	Units
ReflectionSpeedBenchmark.constructorConst	avgt	10	32.392	± 0.473	ns/op
ReflectionSpeedBenchmark.constructorPoly	avgt	10	113.947	± 1.205	ns/op
ReflectionSpeedBenchmark.constructorVar	avgt	10	76.885	± 1.128	ns/op
ReflectionSpeedBenchmark.instanceFieldConst	avgt	10	18.569	± 0.161	ns/op
ReflectionSpeedBenchmark.instanceFieldPoly	avgt	10	98.671	± 2.015	ns/op
ReflectionSpeedBenchmark.instanceFieldVar	avgt	10	54.193	± 3.510	ns/op
ReflectionSpeedBenchmark.instanceMethodConst	avgt	10	33.421	± 0.406	ns/op
ReflectionSpeedBenchmark.instanceMethodPoly	avgt	10	109.129	± 1.959	ns/op
ReflectionSpeedBenchmark.instanceMethodVar	avgt	10	90.420	± 2.187	ns/op
ReflectionSpeedBenchmark.staticFieldConst	avgt	10	19.080	± 0.179	ns/op
ReflectionSpeedBenchmark.staticFieldPoly	avgt	10	92.130	± 2.729	ns/op
ReflectionSpeedBenchmark.staticFieldVar	avgt	10	53.899	± 1.051	ns/op
ReflectionSpeedBenchmark.staticMethodConst	avgt	10	35.907	± 0.456	ns/op
ReflectionSpeedBenchmark.staticMethodPoly	avgt	10	102.895	± 1.604	ns/op
ReflectionSpeedBenchmark.staticMethodVar	avgt	10	82.123	± 0.629	ns/op

### Risks and Assumptions

Code that depends upon highly implementation-specific and undocumented aspects of the existing implementation may be impacted. To mitigate this compatibility risk, as a workaround you can enable the old implementation via `-Djdk.reflect.useDirectMethodHandle=false`.

- Code that inspects the internal generated reflection classes (i.e., subclasses of `MagicAccessorImpl`) will no longer work and must be updated.
- Method-handle invocation may consume more resources than the old core reflection implementation. Such invocation involves calling multiple Java methods to ensure that the declaring class of a member is initialized prior to access, and thus may require more stack space for the necessary execution frames. This may result in a `StackOverflowError` or, if a `StackOverflowError` is thrown while initializing a class, then a `NoClassDefFoundError`.
- We will remove the old core reflection implementation in a future release. The `-Djdk.reflect.useDirectMethodHandle=false` workaround will stop working at that point.