

JEP 445: Unnamed Classes and Instance Main Methods (Preview)

Author	Ron Pressler
Owner	Jim Laskey
Type	Feature
Scope	SE
Status	Closed / Delivered
Release	21
Component	specification / language
Discussion	amber dash dev at openjdk dot org
Effort	5
Relates to	JEP 463: Implicitly Declared Classes and Instance Main Methods (Second Preview)
Reviewed by	Alex Buckley, Brian Goetz
Endorsed by	Brian Goetz
Created	2023/02/13 13:58
Updated	2023/11/09 13:24
Issue	8302326

Summary

Evolve the Java language so that students can write their first programs without needing to understand language features designed for large programs. Far from using a separate dialect of Java, students can write streamlined declarations for single-class programs and then seamlessly expand their programs to use more advanced features as their skills grow. This is a preview language feature.

Goals

- Offer a smooth on-ramp to Java so that educators can introduce programming concepts in a gradual manner.
- Help students to write basic programs in a concise manner and grow their code gracefully as their skills grow.
- Reduce the ceremony of writing simple programs such as scripts and command-line utilities.
- Do not introduce a separate beginner's dialect of Java.
- Do not introduce a separate beginners' toolchain; student programs should be compiled and run with the same tools that compile and run any Java program.

Motivation

Java is a multi-paradigm language that excels for large, complex applications developed and maintained over many years by large teams. It has rich features for data hiding, reuse, access control, namespace management, and modularity which allow components to be cleanly composed while being developed and maintained independently. With these features components can expose well-defined interfaces for their interaction with other components and hide internal implementation details to permit the independent evolution of each. Indeed, the object-oriented paradigm itself is designed for plugging together pieces that interact through well-defined protocols and abstract away implementation details. This composition of large components is called *programming in the large*. Java also offers many constructs useful for *programming in the small* — everything that is internal to a component. In recent years, Java has enhanced both its programming-in-the-large capabilities with modules and its programming-in-the-small capabilities with data-oriented programming.

Java is also, however, intended to be a first programming language. When programmers first start out they do not write large programs, in a team — they write small programs, alone. They have no need for encapsulation and namespaces, useful to separately evolve components written by different people. When teaching programming, instructors start with the basic programming-in-the-small concepts of variables, control flow, and subroutines. At that stage there is no need for the programming-in-the-large concepts of classes, packages, and modules. Making the language more welcoming to newcomers is in the interest of Java veterans but they, too, may find it pleasurable to write simple programs more concisely, without any programming-in-the-large scaffolding.

Consider the classic *Hello, World!* program that is often used as the first program for Java students:

```
public class HelloWorld {
    public static void main(String[] args) {
        System.out.println("Hello, World!");
    }
}
```

There is too much clutter here — too much code, too many constructs, too many constructs — for what the program does.

- The class declaration and the mandatory public access modifier are programming-in-the-large constructs. They are useful when encapsulating a code unit with a well-defined interface to external components, but rather pointless in this little example.
- The `String[] args` parameter also exists to interface the code with an external component, in this case the operating system's shell. It is mysterious and unhelpful here, especially since it is never used.
- The static modifier is part of Java's class-and-object model. For the novice, static is not just mysterious but harmful: To add more methods or fields that `main` can call and use the student must either declare them all as static — thereby propagating an idiom which is neither common nor a good habit — or else confront the difference between static and instance members and learn how to instantiate an object.

The new programmer encounters these concepts at the worst possible time, before they learn about variables and control flow, and when they cannot appreciate the utility of programming-in-the-large constructs for keeping a large program well organized. Educators often offer the admonition, "don't worry about that, you'll understand it later." This is unsatisfying to them and their students alike, and leaves students with the enduring impression that Java is complicated.

The motivation for this JEP is not merely to reduce ceremony but to help programmers that are new to Java, or to programming in general, learn Java in a manner that introduces concepts in the right order: Start with the fundamental programming-in-the-small concepts, then proceed to advanced programming-in-the-large concepts when they are actually beneficial and can be more easily grasped.

We propose to do this not by changing the structure of the Java language — code is still enclosed in methods, which are enclosed in classes, which are enclosed in packages, which are enclosed in modules — but by hiding these details until they are useful in larger programs. We offer an on-ramp, a gradual incline that gracefully merges onto the highway. When students move on to larger programs they need not discard what they learned in the early stages, but rather they see how it all fits within the larger picture.

The changes we offer here are just one step in making Java easier to learn. They do not even address all the speed bumps in the above *Hello, World!* program: The beginner may still be puzzled by the mysterious `System.out.println` incantation, and still needs to import basic utility classes and methods for essential functionality even in first-week programs. We may address these pains in a future JEP.

Description

First, we enhance the protocol by which Java programs are launched to allow *instance main methods*. Such methods are not static, need not be public, and need not have a `String[]` parameter. Then we can simplify the *Hello, World!* program to:

```
class HelloWorld {
    void main() {
        System.out.println("Hello, World!");
    }
}
```

Second, we introduce *unnamed classes* to make the class declaration implicit:

```
void main() {
    System.out.println("Hello, World!");
}
```

This is preview language feature, disabled by default

To try the examples below in JDK 21 you must enable preview features as follows:

- Compile the program with `javac --release 21 --enable-preview Main.java` and run it with `java --enable-preview Main`, or,
- When using the source code launcher, run the program with `java --source 21 --enable-preview Main.java`

The launch protocol

New programmers want to write and run a computer program, but the *Java Language Specification* focuses on defining the core Java unit of the class and the basic compilation unit, namely a source file comprised of a package declaration, followed by some `import` declarations, followed by one or more `class` declarations. All it has to say about a Java *program* is this:

The Java Virtual Machine starts execution by invoking the method `main` of some specified class or interface, passing it a single argument which is an array of strings.

The JLS further says:

The manner in which the initial class or interface is specified to the Java Virtual Machine is beyond the scope of this specification, but it is typical, in host environments that use command lines, for the fully qualified name of the class or interface to be specified as a command line argument and for following command line arguments to be used as strings to be provided as the argument to the method `main`.

The actions of choosing the class containing the main method, assembling its dependencies in the form of a module path or a class path (or both), loading the class, initializing it, and invoking the main method with its arguments constitute the *launch protocol*. In the JDK it is implemented by the *launcher*, i.e., the Java executable.

A flexible launch protocol

We enhance the launch protocol to offer more flexibility in the declaration of a program's entry point and, in particular, to allow *instance* main methods, as follows:

- Allow the main method of a launched class to have public, protected, or default (i.e., package) access.
- If a launched class contains no static main method with a `String[]` parameter but does contain a static main method with no parameters, then invoke that method.
- If a launched class has no static main method but has a non-private zero-parameter constructor (i.e., of public, protected, or package access), and a non-private instance main method, then construct an instance of the class. If the class has an instance main method with a `String[]` parameter then invoke that method; otherwise, invoke the instance main method with no parameters.

These changes allow us to write *Hello, World!* with no access modifiers, no static modifiers, and no `String[]` parameter, so the introduction of these constructs can be postponed until they are needed:

```
class HelloWorld {
    void main() {
        System.out.println("Hello, World!");
    }
}
```

Selecting a main method

When launching a class, the launch protocol chooses the first of the following methods to invoke:

- A static `void main(String[] args)` method of non-private access (i.e., public, protected or package) declared in the launched class,
- A static `void main()` method of non-private access declared in the launched class,
- A `void main(String[] args)` instance method of non-private access declared in the launched class or inherited from a superclass, or, finally,
- A `void main()` instance method of non-private access declared in the launched class or inherited from a superclass.

Note that this is a change of behavior: If the launched class declares an instance `main`, that method will be invoked rather than an inherited "traditional" `public static void main(String[] args)` declared in a superclass. Therefore, if the launched class inherits a "traditional" main method but another method (i.e. an instance `main`) is selected, the JVM will issue a warning to the standard error at runtime.

If the selected `main` is an instance method and is a member of an inner class, the program will fail to launch.

Unnamed classes

In the Java language, every class resides in a package and every package resides in a module. These namespacing and encapsulation constructs apply to all code, but small programs that do not need them can omit them. A program that does not need class namespaces can omit the package statement, making its classes implicit members of the unnamed package; classes in the unnamed package cannot be referenced explicitly by classes in named packages. A program that does not need to encapsulate its packages can omit the module declaration, making its packages implicit members of the unnamed module; packages in the unnamed module cannot be referenced explicitly by packages in named modules.

Before classes serve their main purpose as templates for the construction of objects, they serve only as namespaces for methods and fields. We should not require students to confront the concept of classes before they are comfortable with the more basic building blocks of variables, control flow, and subroutines, before they embark on learning object orientation, and when they are still writing simple, single-file programs. Even though every method resides in a class, we can stop requiring explicit class declarations for code that does not need it — just as we do not require explicit package or module declarations for code that does not need them.

Henceforth, when the Java compiler encounters a source file with a method that is not enclosed in a class declaration it will implicitly consider such methods, as well as any unenclosed fields and any classes declared in the file, to be members of an *unnamed top-level class*.

An unnamed class is always a member of the unnamed package. It is also `final` and cannot implement any interface nor extend any class other than `Object`. An unnamed class cannot be referenced by name, so there can be no `method references` to its static methods; the `this` keyword can still be used, however, and so can `method references` to instance methods.

No code can refer to an unnamed class by name, so instances of an unnamed class cannot be constructed directly. Such a class is useful only as a standalone program or as an entry point to a program. Therefore, an unnamed class must have a `main` method that can be launched as described above. This requirement is enforced by the Java compiler.

An unnamed class resides in the unnamed package, and the unnamed package resides in the unnamed module. While there can be only one unnamed package (barring multiple class loaders) and only one unnamed module, there can be multiple unnamed classes in the unnamed module. Every unnamed class contains a `main` method and so represents a program, thus multiple such classes in the unnamed package represent multiple programs.

An unnamed class is almost exactly like an explicitly declared class. Its members can have the same modifiers (e.g., `private` and `static`) and the modifiers have the same defaults (e.g., package access and instance membership). One key difference is that while an unnamed class has a default zero-parameter constructor, it can have no other constructor.

With these changes we can now write *Hello, World!* as:

```
void main() {
    System.out.println("Hello, World!");
}
```

Top-level members are interpreted as members of the unnamed class, so we can also write the program as:

```
String greeting() { return "Hello, World!"; }

void main() {
    System.out.println(greeting());
}
```

or, using a field, as:

```
String greeting = "Hello, World!";

void main() {
    System.out.println(greeting);
}
```

If an unnamed class has an instance main method rather than a static main method then launching it is equivalent to the following, which employs the existing *anonymous class declaration* construct:

```
new Object() {
    // the unnamed class's body
}.main();
```

A source file named `HelloWorld.java` containing an unnamed class can be launched with the source-code launcher, like so:

```
$ java HelloWorld.java
```

The Java compiler will compile that file to the launchable class file `HelloWorld.class`. In this case the compiler chooses `HelloWorld` for the class name as an implementation detail, but that name still cannot be used directly in Java source code.

The javadoc tool will fail when asked to generate API documentation for a Java file with an unnamed class, as unnamed classes do not define any API accessible from other classes. This behavior may change in a future release.

The `Class.isSynthetic` method returns `true` for an unnamed class.

Growing a program

A *Hello, World!* program written as an unnamed class is much more focused on what the program actually does, omitting concepts and constructs it does not need. Even so, all members are interpreted just as they are in an ordinary class. To evolve an unnamed class into an ordinary class, all we need to do is wrap its declaration, excluding `import` statements, inside an explicit `class` declaration.

Eliminating the main method altogether may seem like the natural next step, but it would work against the goal of gracefully evolving a first Java program to a larger one and would impose some non-obvious restrictions (see below). Dropping the `void` modifier would similarly create a distinct Java dialect.

- Alternatives**
 - Use JShell for introductory programming* — A JShell session is not a program but a sequence of code snippets. Declarations typed into `jshell` are implicitly viewed as static members of some unspecified class, with some unspecified access level, and statements execute in a context in which all previous declarations are in scope.
 - This is convenient for experimentation — which is the primary use case for JShell — but not a good model for learning to write Java programs. Evolving a batch of working declarations in JShell into a real Java program leads to a non-idiomatic style of code because it declares each method, class, and variable as `static`. JShell is a great tool for exploration and debugging, but it is not the on-ramp programming model we are looking for.
 - Interpret code units as static members* — Methods and fields are non-static by default. Interpreting top-level members in an unnamed class as static would change the meaning of the top-levels in such a class — introducing, in effect, a distinct Java dialect. To preserve the meaning of such members when we evolve an unnamed class into an ordinary class we would have to add explicit static modifiers. This is not what we want as we scale up from a handful of methods to a simple class. We want to start using classes as classes, not as containers of static members.
 - Interpret code units as locals* — We can already declare local variables with `local` methods. Assume that we could also declare local methods, i.e., methods within other methods. Then we could interpret the body of a simple program as the body of a main method, with variables interpreted as local variables rather than fields, and methods interpreted as local methods rather than class members. This would allow us to eschew the main method altogether and write top-level statements.

The problem with this approach is that, in Java, locals behave differently from fields, and in a more restricted way to boot: Locals can only be accessed from inside lambda bodies or inner classes when they are *effectively final*. The proposed design allows us to separate locals and fields in the same manner as is always done in Java. The burden of writing a main method is not onerous, even for new students.

- Introduce package-level methods and fields* — A user experience similar to that shown above could be achieved by allowing package-level methods and fields to be declared in a file without an explicit package or class declaration. However, such a feature would have a far wider impact on how Java code is written in general.